

Redis data persistency on Uberspace

July 31, 2015 | 0 Comments | Tags: technology, berichtsheft, configuration, database, tutorial, Redis, Uberspace

[Uberspace](#) provides a nice [german introduction](#) on how to use Redis on their servers and also kindly points out that Redis resides in memory and will (when used with default settings) lose its state (speak: all your data **yikes**) when there is a crash or restart. So let us fix that:

Setting up Redis

Just a quick rehearsal of their tutorial. Redis is being set up in two steps basically:

```
test -d ~/service || uberspace-setup-svscan
uberspace-setup-redis
```

Enabling persistency

There are several persistency models you can read up on here:. But I would recommend sticking with the most common way shown here, which strikes a good balance between performance and data safety.^[1]

```
nano ~/.redis/conf
```

And then add the following lines:

```
dir /home/YOUR_UBERSPACE_USER_NAME/.redis/
appendonly yes
appendfsync everysec
```

```
# Restart redis
svc -du ~/service/redis
```

Verifying persistency

Now run the tool that will put data into Redis and make it do so. You can then check ① both that the persistency file grew to a size greater than 0 and ② also restart Redis and afterwards check whether your tool still sees the previously entered data.

```
# ① Check that the persistency file holds content
du -sch ~/.redis/appendonly.aof
```

```
# ② Restart Redis. This would kill all the contained data
# if the persistency setting would not work. Then check if
# the data is still available in your Redis using tool
svc -du ~/service/redis
```

Sources

1. [Redis persistence demystified \(oldblog.antirez.com\)](http://oldblog.antirez.com)